


Qwen3-8B · Attribution efficiency

Nominal 10-token input · Updated DeltaTrace and FlashTrace curves



DT / FT: same C550, complete warm API calls; cold costs recorded separately. [ref]: released 6/8-GPU results.
OOM / failed cells (no latency inferred): FT multi-hop: 5000, 10000; FlashTrace: 5000, 10000 tokens.